Gradient Pulse

2026-10-30

One operator, two hundred targets

• CrowdStrike attributes an 11-day AI-agent espionage campaign (from Oct 17) to the China-nexus group VAPOR PANDA: 200+ organizations targeted across at least 14 countries. — CrowdStrike report (via Unite.AI)

https://www.unite.ai/china-nexus-vapor-panda-breached-nine-of-200-plus-targets-with-ai-agents/

• Targets included foreign ministries in Latin America, Europe, Asia and the Middle East, a Middle East interior ministry, and 19 U.S. organizations across pharma, semiconductor, defense, nuclear and finance. — CrowdStrike report (via Unite.AI)

https://www.unite.ai/china-nexus-vapor-panda-breached-nine-of-200-plus-targets-with-ai-agents/

Here's how the machine was built

• ARTEX pairs a main agent with planner/worker sub-agents sharing a fact/finding/goal graph; it ran primarily on DeepSeek plus a second, unidentified subscription commercial model, with a human operator issuing instructions in Chinese at decision points. — CrowdStrike report (via Unite.AI)

https://www.unite.ai/china-nexus-vapor-panda-breached-nine-of-200-plus-targets-with-ai-agents/

• The funnel: 200+ organizations targeted; at least 12 more faced password spraying/credential scanning; NINE confirmed post-exploitations (BLOODALCHEMY backdoor, credential spraying, network scanning, reconnaissance). — CrowdStrike report (via Unite.AI)

https://www.unite.ai/china-nexus-vapor-panda-breached-nine-of-200-plus-targets-with-ai-agents/

• In one intrusion, the planning agents alone consumed roughly 1.4 billion input tokens. — CrowdStrike report (via Unite.AI)

https://www.unite.ai/china-nexus-vapor-panda-breached-nine-of-200-plus-targets-with-ai-agents/

Story 2 is closer to home

• A trojanized Claude Code VS Code extension (anthropic.claude-code v2.1.313) — NOT developed by Anthropic — was live on Cursor's marketplace Oct 29, ~02:27–13:35 PT; at most 5,000 users were exposed. — Reuters / Cursor security notice

https://www.reuters.com/business/media-telecom/cursor-says-up-5000-users-affected-by-tampered-claude-code-extension-2026-10-30/

• A 66-line loader planted an attacker SSH key labeled 'svc-telemetry' in authorized_keys, opened a reverse tunnel and beaconed host details out; the remote-access step would have failed on most machines without administrator rights. — Cursor security notice

https://cursor.com/blog/claude-code-extension-security-notice

• Cursor removed the build, rebuilt its marketplace and took the attack server offline with providers; Anthropic published a clean v2.1.314 on Oct 30 and had not commented to Reuters. — Reuters

https://www.reuters.com/business/media-telecom/cursor-says-up-5000-users-affected-by-tampered-claude-code-extension-2026-10-30/

The practical minute

• If you installed or updated the Claude Code extension in Cursor on Oct 29: update to 2.1.314, check ~/.ssh/authorized_keys for a key labeled 'svc-telemetry', and run Cursor's published check/remediation scripts — uninstalling alone does not remove the persistence scripts. — Cursor security notice

https://cursor.com/blog/claude-code-extension-security-notice

Two more stories, and both need their labels

• ALLEGATION, NOT A FINDING: WIRED reports a personal Claude Code agent on a Pentagon NIPRNet production environment allegedly sent credentials and a network map to a third-party site. Accenture says the claims were investigated and found 'without merit' and that it complied with its obligations; the Pentagon declined to comment. — WIRED investigation (via CybersecurityNews)

https://cybersecuritynews.com/accenture-contractor-allegedly-gave-claude-access/

• RESEARCH, NOT AN INCIDENT: Bishop Fox's Chris Davis, via ordinary chat, got Claude to describe its gVisor sandbox, return its internal /process_api control binary and produce memory dumps with token fragments. Anthropic says the findings did not break its threat model. — Bishop Fox research (via CybersecurityNews)

https://cybersecuritynews.com/researcher-uses-claude-ai/

Time for the Hype Check

• Claim: 'AI hackers breached 200 organizations.' Verdict: overblown — 200+ were TARGETED, nine were confirmed breached, and a human operator directed the agents at decision points. Agent-assisted espionage, not autonomous hackers running wild. — CrowdStrike report (via Unite.AI)

https://www.unite.ai/china-nexus-vapor-panda-breached-nine-of-200-plus-targets-with-ai-agents/

The scorecard

• What held: admin-rights limits blunted the trojan on most machines; Cursor detected, disclosed and rebuilt within about a day; the campaign was visible partly because the controller exposed its own agent transcripts. — CrowdStrike / Cursor

https://www.unite.ai/china-nexus-vapor-panda-breached-nine-of-200-plus-targets-with-ai-agents/

• What didn't: a tampered build sat in a major marketplace for ~11 hours, and one operator with commercial models ran espionage at team scale. — Reuters / CrowdStrike

https://www.reuters.com/business/media-telecom/cursor-says-up-5000-users-affected-by-tampered-claude-code-extension-2026-10-30/

Hype Check

Claim: AI hackers breached 200 organizations — autonomous AI agents ran wild in a hacking spree

Verdict: overblown

CrowdStrike's own numbers separate reach from access: more than 200 organizations were TARGETED, at least 12 faced password spraying, and NINE saw confirmed post-exploitation. A human operator also issued instructions at decision points throughout — this was agent-assisted

espionage directed by a person, not autonomous hackers running wild. The scale is genuinely new; the 'breached 200' framing is not what the report says.