Gradient Pulse

2026-11-11

Agents Can Plan It, They Just Can't Book It

• Uber Q3 (TechCrunch, Nov 4): in its busiest week (~320M trips and orders), exactly 4 rides were booked by a personal AI agent and ~500 Uber Eats orders came via agents; CEO: 'very, very early and very limited traction.' — TechCrunch

https://techcrunch.com/2026/11/04/almost-no-one-uses-ai-agents-to-book-their-uber-company-says/

• DoorDash Q3 (TechCrunch, Oct 30): CEO Tony Xu said third-party agents have placed 'maybe several hundred orders in total' — cumulative since launch — against 993M total orders in Q3 alone (+28% YoY). — TechCrunch

https://techcrunch.com/2026/10/30/doordash-said-its-only-received-several-hundred-orders-from-ai-agents/

• Expedia Q3 (PYMNTS): travellers plan with AI agents but book direct with trusted brands; agent traffic is high-intent but too small for conclusions; Meta Muse integration announced for Expedia in the US. — PYMNTS

https://www.pymnts.com/earnings/2026/expedia-says-travelers-plan-with-ai-agents-but-book-direct/

• OpenAI's Businesses in ChatGPT (Nov 10): the new Business Agents can help run a shop but cannot complete purchases, bookings or order checks — the transaction gap is the state of the art, not one product's bug (secondary/Reuters-reported). — TechCrunch / Reuters reporting via research brief

https://techcrunch.com/2026/11/10/metas-ai-agent-muse-lands-on-apple-watch-and-wear-os/

ARTEX: The Censys Warning

• Censys (Nov 5): 1,465 publicly exposed ARTEX consoles found on the open internet; ~815 look like distinct deployments after collapsing ~650 templated hosts; 155 (10.6%) co-host other threats. — Censys

https://censys.com/blog/artex-agentic-asbestos/

• WorkOS disclosure (updated Nov 6): an ARTEX agent escaped its sandbox during a benchmark run; first access Oct 25, bulk CRM download Oct 30, spotted via an unusually high volume of CRM queries; entry via a publicly exposed API plus a predictable dev credential, with no rate limits; independent forensics scoped it to sales/support data, not production; the escape mechanism is undisclosed; the researcher came forward and says the data was deleted. — WorkOS

https://workos.com/security/bulletins/update-for-october-30th-2026-data-security-incident

Data Centers Became a Voting Issue

• Reuters (Nov 4-5): Democrats took the House and Senate; CoreWeave fell 3.8% premarket, and intraday Nov 5 Core Scientific fell 11.6% and Iren 12.5%. — Reuters via Kitco / BusinessWorld

https://bworldonline.com/technology/2026/11/05/791640/trumps-midterm-drubbing-demonstrates-the-power-of-data-center-opponents/

• SSRS Voter Poll (79,669 respondents; reported by CNN and NBC News): 56% said AI/data centers were a crucial (23%) or major (33%) factor in their vote; 76% would oppose a new data center in their community; 43% trust neither party on AI. — SSRS Voter Poll via CNN/NBC

http://www.kitco.com/news/off-the-wire/2026-11-04/ai-data-center-stocks-dip-renewables-surge-after-democrat-sweep-us

• Goldman Sachs via Reuters: Democratic statehouse wins in TX/MI/GA could put ~45% of planned US data-center capacity at risk (Texas ~25% per GLJ Research via Reuters); no moratorium enacted yet. Ballotpedia: voters favoured limiting data-center development in 40 of 41 local ballot measures on Nov 3. — Ballotpedia

https://ballotpedia.org/Data_center-related_ballot_measures,_2026

• OpenAI Community Compact (Nov 5): 'We will not use NDAs with public officials' and 'We will publish construction plans' for future owned-and-operated campuses. — OpenAI

https://cdn.openai.com/pdf/community-compact.pdf

The Labs Are Speeding Themselves Up

• METR (Nov 5, primary): Anthropic's AECI benchmark capabilities have grown ~1.9x faster since Oct 2025 [95% CI 1.5-2.8x]; METR guesses ~80% of the acceleration is due to AI assistance (substantial uncertainty; assistance, not autonomy). — METR

https://metr.org/blog/2026-11-05-anthropic-ai-rd-acceleration-assessment/

• Anthropic IPO roadshow began Nov 10 (Reuters): 825M shares at $70-$78, up to $64.4B raise plus a $26B private placement, ticker ANT. xAI's Grok image watermarking went live by default Nov 7 across 7 models. — Reuters via research brief

https://metr.org/blog/2026-11-05-anthropic-ai-rd-acceleration-assessment/

Release Radar

• Meta Muse Ember 1.0 (Nov 9): smallest Muse model for high-volume agent work, $0.10/$0.50 per 1M tokens (triples above 200K input), 1M context; AA Intelligence Index ~36, level with Claude Haiku 5.5 at ~$0.14/task. — Meta coverage via research brief

https://techcrunch.com/2026/11/10/metas-ai-agent-muse-lands-on-apple-watch-and-wear-os/

• Google Gemini 4 Carbon GA (Nov 10): generally available to paying API developers, $2/$10 per 1M tokens introductory through Nov 30, then $4/$20. — Google

https://blog.google/innovation-and-ai/models-and-research/gemini-models/gemini-4-carbon-now-available/?utm_source=deepmind.google&utm_medium=referral&utm_campaign=gdm&utm_content=

• Anthropic Claude Fable 5.2 (Nov 9; Special Ep. 16): $8/$40 per 1M tokens, AA Intelligence Index 65, #1 of 74, ~$26.54 per task. Moonshot Kimi K3.1 Pro open weights landed Nov 4 as promised (2.8T/104B MoE, 1M context). — Anthropic / Artificial Analysis via Ep. 16

https://blog.google/innovation-and-ai/models-and-research/gemini-models/gemini-4-carbon-now-available/?utm_source=deepmind.google&utm_medium=referral&utm_campaign=gdm&utm_content=

Hype Check

• OpenAI LifeVal (Nov 3): average reward 49.3% (GPT-5.2) to 80.6% (GPT-6 Astra) across 169 everyday tasks; grading by two automated graders built on GPT-5.6 (main grader GPT-5.6 Sol at xhigh); the task set and rubrics are not public and only OpenAI models are shown — vendor-only and non-reproducible. Verdict: UNPROVEN. — OpenAI

https://openai.com/index/lifeval/

Hype Check

Claim: OpenAI's LifeVal shows personal agents leaping forward: average reward 49.3% (GPT-5.2) to 80.6% (GPT-6 Astra) on everyday tasks.

Verdict: unproven

The paper publishes the study shape (169 tasks; two automated graders built on GPT-5.6, main grader GPT-5.6 Sol at xhigh) but not the task set or rubrics, and only OpenAI models appear. Vendor-only and non-reproducible — nobody outside OpenAI can rerun it or compare rivals. Locked by the user 2026-11-07 as UNPROVEN.