Gradient Pulse

2026-10-28

Two weeks after Argon

• Google DeepMind announced Gemini 4 Carbon today, October 28 — its second frontier model in two weeks, after Gemini 4 Argon on October 14 — Google DeepMind

https://blog.google/innovation-and-ai/models-and-research/gemini-models/introducing-gemini-4-carbon/

• Google calls Carbon 'our strongest model with state-of-the-art cybersecurity and long horizon capabilities'; the announcement was authored by DeepMind SVP Koray Kavukcuoglu — Google DeepMind

https://blog.google/innovation-and-ai/models-and-research/gemini-models/introducing-gemini-4-carbon/

• Reuters confirms the launch, the cybersecurity focus and the two-weeks-after-Argon cadence — Reuters

https://www.reuters.com/technology/google-launches-gemini-4-carbon-ai-model-with-cybersecurity-focus-2026-10-28/

Number one on the index

• Carbon ranks #1 on the updated Artificial Analysis Intelligence Index v5.0 at 59 — an independent scorecard, not Google's own — Artificial Analysis via Google

https://blog.google/innovation-and-ai/models-and-research/gemini-models/introducing-gemini-4-carbon/

• Claude Opus 5.5 scores 58 and GPT-6 Astra 57 on the same v5.0 index; Gemini 4 Argon scores 50 — Artificial Analysis

https://artificialanalysis.ai/

• The rebuilt v5.0 index puts most of its weight on private test sets the labs cannot study in advance — Artificial Analysis via press

https://www.reuters.com/technology/google-launches-gemini-4-carbon-ai-model-with-cybersecurity-focus-2026-10-28/

The cyber numbers

• SEC-bench Pro 88.2% for finding vulnerabilities in real-world software; ExploitGym 54.8% for proving flaws can be exploited; CWE-bench 71.0% for writing the patch — all Google-run results — Google DeepMind

https://blog.google/innovation-and-ai/models-and-research/gemini-models/introducing-gemini-4-carbon/

• Google states the three cyber benchmarks were 'evaluated without guardrails' — safety filters off, so the scores are capability ceilings, not shipped-product behavior — Google DeepMind

https://blog.google/innovation-and-ai/models-and-research/gemini-models/introducing-gemini-4-carbon/

• On internal benchmarks, Carbon generates secure code at more than twice Argon's success rate — Google DeepMind

https://blog.google/innovation-and-ai/models-and-research/gemini-models/introducing-gemini-4-carbon/

Built to work for days

• Google's internal teams have run Carbon for days at a time on multi-step engineering work — root-causing and patching subtle bugs across large codebases — Google DeepMind

https://blog.google/innovation-and-ai/models-and-research/gemini-models/introducing-gemini-4-carbon/

• Internal testing: unsupported claims cut roughly in half versus Argon; high-severity misalignment alerts in live coding down almost 80%; successful cheating on Terminal-Bench 4.0 reduced to 0.0% — Google DeepMind

https://blog.google/innovation-and-ai/models-and-research/gemini-models/introducing-gemini-4-carbon/

• Gray Swan indirect prompt-injection benchmark: 0.0% attack success for Carbon, versus 0.7% for Argon — all internal Google figures — Google DeepMind

https://blog.google/innovation-and-ai/models-and-research/gemini-models/introducing-gemini-4-carbon/

Almost nobody can use it

• Rollout is restricted: starting this week, a subset of cyber defenders gets Carbon, phasing to more trusted partners through Google's Fairwind Program; Google is in the US government's voluntary pre-release testing process — Google DeepMind

https://blog.google/innovation-and-ai/models-and-research/gemini-models/introducing-gemini-4-carbon/

• No date for developers, businesses or consumers — Google says 'as soon as possible' — Reuters

https://www.reuters.com/technology/google-launches-gemini-4-carbon-ai-model-with-cybersecurity-focus-2026-10-28/

• Pricing is set anyway: $2 per million input tokens and $10 per million output at introductory rates matching Argon, cached input 95% off; the standard rate card is $4/$20 — Google DeepMind / Artificial Analysis

https://blog.google/innovation-and-ai/models-and-research/gemini-models/introducing-gemini-4-carbon/

Hype Check: today's claim

• Claim: Google's cybersecurity benchmark sweep proves Carbon is state-of-the-art at finding, exploiting and patching vulnerabilities — Google DeepMind

https://blog.google/innovation-and-ai/models-and-research/gemini-models/introducing-gemini-4-carbon/

• Every cyber number is Google-run and Google-reported, measured without guardrails, with no independent rerun yet; the independently measured leg — the Artificial Analysis Index #1 — measures general capability, not cyber skill — Gradient Pulse analysis

https://www.reuters.com/technology/google-launches-gemini-4-carbon-ai-model-with-cybersecurity-focus-2026-10-28/

The scorecard

• Carbon joins the release radar at $2/$10 introductory pricing, closed weights, restricted access; Argon remains the Gemini 4 model anyone can actually buy — Gradient Pulse

https://www.agilefactornews.com/radar/ai/?episode=special-gemini-4-carbon-2026-10-28

• What to watch: a general-availability date, and the first independent rerun of the cyber benchmarks with guardrails on — Gradient Pulse

https://www.agilefactornews.com/radar/ai/?episode=special-gemini-4-carbon-2026-10-28

Hype Check

Claim: Google's cybersecurity benchmark sweep — SEC-bench Pro 88.2%, ExploitGym 54.8%, CWE-bench 71.0% — proves Gemini 4 Carbon is state-of-the-art at finding, exploiting and patching vulnerabilities

Verdict: unproven

Every cyber number is Google-run and Google-reported, measured without guardrails, with no independent rerun yet. The independently measured leg — #1 on the Artificial Analysis Intelligence Index v5.0 at 59 — measures general capability, not cyber skill. The scores may hold; today they are a vendor scorecard.