

GRADIENT PULSE SPECIAL

"I Tested Hark AI's Pro: 6 Real Tasks Tried, Timed"

Special Edition · October 7, 2026 · Agile Factor

A chaptered, first-person hands-on review. Six real tasks, timed. Verdict: split — the hands are confirmed, the brain needs a second opinion.

Format: chaptered hands-on review | Target: 12–15 min (~2,150 words)

COLD OPEN (~30s)

■ [VISUAL: Hark logo / headline cards — "97.7", "\$700M", "best internet-use model ever built"]

An AI startup you've probably never used just claimed it built the best internet-using AI on Earth. Ninety-seven point seven on the benchmark that measures it — that's a test that scores how well an AI can actually operate websites, not just chat about them. Cheaper than the frontier models by a factor of ten. And seven hundred million dollars in fresh funding says they're serious.

I got invite access to Hark Pro. And instead of reading you the press release, I gave it jobs. Dinner. Flowers. Flights. A captcha. A memory test. Told it to time every single one. Here's what actually happened.

CHAPTER 1 — SETUP: WHAT HARK IS CLAIMING (~2 min)

■ [VISUAL: RuntimeWire / TechCrunch / BusinessWire headline cards]

Hark is a consumer AI platform that launched paid plans this week. Back in May they raised a reported seven hundred million dollar Series A at a six billion dollar valuation — that's serious money for a company most people hadn't heard of. Their first hundred thousand signups were promised the paid plan free.

The product has two halves. There's the browser agent, called Handoff — that's the part that goes and does things on websites for you. And there's the platform around it: personalized panels, memory, scheduled tasks.

The claim we're testing today is about Handoff. Hark says it scored 97.7 on Online-Mind2Web — the web-use benchmark — against 92.8 for GPT-5.4, 84.1 for Claude Opus 4.8, 69 for Gemini 2.5 Pro. And they price it at eighteen cents per million input tokens — they say under a tenth of frontier pricing. One caveat before we start, flagged by the outlet Wellfunded: those comparisons leave out the newest OpenAI and Anthropic models. Keep that in your pocket.

My method: seven prompts, each instructed to time itself, each told to stop before spending any money. No payment method connected. Let's go.

CHAPTER 2 — DINNER: THE ORDER TEST (~1.5 min)

■ [VISUAL: test1-doordash-login.png, then test1-doordash-order-card.png]

Task one: order dinner for two on DoorDash, delivery around seven, stop at checkout.

First pleasant surprise: logging into DoorDash happened right in the chat through an encrypted vault. No password pasted into a sketchy form. Clean.

Then it asked me to confirm my city and zip — and caught that the street name on file was wrong, suggesting the correction itself. That's not clicking buttons. That's noticing.

Eleven minutes total, about nine of actual browser work. Final card: a local Italian place, combo plus lasagna, \$52.31, held at checkout before their cutoff. Nothing charged.

Verdict on task one: complete. And the standout wasn't the ordering — it was the catching.

CHAPTER 3 — FLOWERS: THE FUZZY TEST (~1.5 min)

■ [VISUAL: [test2-florist-browser.png](#), [test2-florist-cart.png](#), [test2-florist-result-card.png](#)]

Task two was deliberately vague: build a bouquet with "some of the florist's choice flowers," under sixty dollars.

This is the test most agents fail, because there's nothing to click — there are decisions to make. Hark made four: it picked the top-rated local florist in my hometown, read "under sixty" as all-in including delivery, chose a forty-dollar Designer's Choice bouquet — which is literally the florist's choice — in bright mixed seasonal flowers, for delivery tomorrow. Then it reported every one of those calls and offered to change any of them before anything moved forward.

Five minutes flat. And note where this happened: not Amazon, not a big platform — a small local florist's website. The "unfamiliar websites" claim holding up in the wild.

Two for two. And a pattern emerging: the impressive moments aren't navigation. They're judgment.

CHAPTER 4 — PRICE CHECK: THE OBSTACLE COURSE (~1.5 min)

■ [VISUAL: [test4-price-table.png](#)]

Task three: cheapest Sony WH-1000XM5 headphones across three retailers. I told it to count every obstacle — pop-ups, banners, walls.

Three minutes. Fastest task yet. One obstacle the whole run: an Amazon captcha page. It got past a captcha — the thing explicitly built to stop exactly this kind of agent — and kept going.

\$198 everywhere, about half off. But here's the moment that matters: Walmart's \$198 wasn't Walmart. It was a third-party seller called Focus Camera. Hark flagged that and recommended Best Buy instead — same price, sold directly, in stock. It rendered the whole thing as a clean comparison table: retailer, price, sold-by, obstacles.

It didn't just report prices. It made a buying call, with reasoning. Three for three.

CHAPTER 5 — FLIGHTS: THE STUMBLE (~2 min)

■ [VISUAL: [test5-flights-browser.png](#), then [test5-flights-table.png](#)]

Task four: cheapest round-trip to Denver next month, top three in a table. Research only, no booking.

Eight minutes. The table was right: three Frontier nonstops, \$43 to \$46. It even recommended one — November 4th to 11th at \$43 — and reasoned about total cost, noting Frontier charges extra for carry-ons.

And then it said something false. Confidently. It claimed Southwest's \$131 fare includes two free checked bags. That policy is gone — Southwest ended it. I corrected it, and to its credit, it said "my mistake," and revised the recommendation on the spot: Southwest and American both include a carry-on at that price, so they're the better pick if you're bringing more than a personal item.

This is the most important result of the day. The browsing was flawless. The failure was in what it *knew* — stale knowledge, delivered with total confidence. The hands worked. The brain was stale. Remember that; it's coming back in the Hype Check.

CHAPTER 6 — MEMORY: THE PAYOFF (~2 min)

■ [VISUAL: [test6-memory-chat-redacted.png](#), [test6-memory-multitab.png](#), [test6-memory-results-redacted.png](#) — REDACTED VERSIONS ONLY]

Earlier in the day I'd told it: remember this — I always want aisle seats, and I prefer morning departures. Then, separately: find me a round-trip flight to Chicago next week. No mention of the preferences. The test is whether it applies them unprompted.

It did — and then some. It picked morning departures. It flagged that Frontier's cheap fare charges extra for the aisle seat and the carry-on. And it recommended American at \$546 as, quote, "worth the extra \$115" for a morning flight home with a carry-on included. That's not recall. That's reasoning with a preference — weighing my aisle-seat habit against a hundred-and-fifteen-dollar gap and telling me when the upgrade earns it.

It also remembered my delivery address from the DoorDash test without being told, and paired the DFW flight with a drive-time departure plan — about four hours from my hometown, so leave by 3:30 AM. All while juggling multiple browser tabs.

This is the memory feature doing exactly what the marketing says it does. Strongest result of the day.

CHAPTER 7 — THE METER: WHAT IT COST (~1 min)

Last task wasn't a task — it was an audit. After six real browser jobs, I'd used one percent of my daily allowance. Zero percent of the monthly pool. Daily resets tomorrow afternoon, monthly in November, and if I ever blow past the allowance it auto-tops-up instead of stopping mid-task.

Two honest limitations: there's no per-task cost breakdown — just overall percentages under billing. And the tier names are a maze: I'm on something called Pro³, there's a Pro² at around twenty a month, and Hark says the free tier keeps every feature — browsers, panels, memory, all of it — just with less usage.

Six tasks. One percent. Whatever the absolute numbers are, the allowance is generous.

CHAPTER 8 — SIDEBAR: THE PANEL THING (~1.5 min)

■ [VISUAL: [panel-running-weather.png](#)]

Between tests I tried something that isn't a browser task at all. I told it: "I'm an avid runner. Build a weather panel for my location telling me how ideal the weather is for running."

It built me a persistent dashboard panel. "Fair for running right now — warm with low humidity." Eighty-two degrees, humidity, dew point, hourly outlook. It's still there.

This reframes the whole video. The browser agent is the demo. But panels like this are the actual product thesis: not an agent that does tasks, but a platform that molds itself around your life — that knows you're a runner, remembers the aisle seat, catches the wrong street name. Hark isn't selling clicks. It's selling a dashboard for your life that happens to be able to click.

SITUATION ROOM SPOT (~30s)

■ [VISUAL: Situation Room branding, GP Ledger / Arena cards]

[YOUTUBE STUDIO CARD: link card → Situation Room signup CTA, timed to this segment]

If you liked watching us stress-test Hark's claims instead of reading the press release — that's what the Situation Room is. Every week: the full Gradient Pulse ledger, every AI claim on the record; Agent Arena premium, with round-by-round evidence and rematches when the dust settles; and Hype Check follow-ups on the claims we called. Five dollars a month, seven-day trial. Link's below — and tap the card on screen to join.

HYPE CHECK (~1.5 min)

■ [VISUAL: Hype Check card — claim, evidence, verdict: SPLIT]

The claim: "Hark Handoff is the best internet-use model ever built." Ninety-seven point seven on the web-use benchmark, under a tenth of the token price.

What we verified: six tasks, six completions. A captcha cleared. Judgment calls on every fuzzy instruction. Memory that reasons, not just recalls. The web-use layer — the thing the benchmark actually measures — held up impressively.

What we didn't verify: we didn't reproduce the benchmark, and "best ever built" is still their superlative, against a comparison set that skips the newest models. And the one failure we found — the Southwest baggage fiction — sits entirely outside what that benchmark measures. Online-Mind2Web scores navigation. It doesn't score knowing things.

So the verdict is split. The hands: confirmed. The brain: get a second opinion. And the superlative stays on the vendor's tab until somebody independent reproduces it.

CLOSER (~1 min)

Six tasks. Three to eleven minutes each. Every one of them slower than doing it myself — and not one of them cost me a minute of attention. That's the trade, stated honestly: Hark doesn't save you time. It saves you attention.

And the real bet here isn't the browser agent. It's the platform underneath — the thing building running-weather panels, remembering aisle seats, catching wrong street names. An agent does your tasks. A platform learns your life. That's the more interesting story, and it's the one I'll be watching.

Gradient Pulse. We verify the claims so you don't have to. Subscribe, and we'll see you in the next one.