

Gradient Pulse — AI News, Explained

2026-09-21 · Special edition · Issued 14:25 America/Chicago

What shipped

- Grok 4.7 shipped September 21 as xAI's new flagship, 40 days after Grok 4.6; live in the Grok app, Cursor, Grok Build, the xAI API, GitHub Copilot, and major model routers. — [xAI / AI Release Tracker](#)
- Claimed 2.1T parameters (+40% over 4.6) — xAI publishes no architecture details, model card, or system card, so the count cannot be independently verified. — [xAI \(Musk's figure; no spec sheet\)](#)
- 500K token context window (unchanged from 4.6); pricing \$2/\$6 per 1M tokens, unchanged — doubling above 200K tokens, cached input \$0.50/\$1.00. — [AI Release Tracker](#)
- Supplemental training on SpaceX engineering data — Starlink telemetry, manufacturing records, failure logs — pitched as better reasoning about hardware/physical systems; no independent verification possible. — [xAI \(company claim\)](#)

Hype Check: the benchmark gap

- Self-reported: DeepSWE 71.0% (run at High, not xHigh effort), CursorBench 46.3%, Harvey Legal Agent Benchmark 19.6%, Terminal-Bench 4.0 38.0%, EEBench 64.0% — all self-run, no third-party rerun attached. — [xAI launch materials](#)
- Rival-model figures in xAI's tables are xAI's own runs, not the labs' published results; the X post and the launch page used different comparison sets (launch page swapped in GPT-6 Astra). — [xAI launch materials](#)
- Intelligence Index v4.3.2: Grok 4.7 = 46 vs 53 for GPT-6 and 53 for Claude Fable 5.1 — mid-pack, seven points back, +2 over Grok 4.6. — [Artificial Analysis \(independent\)](#)
- Terminal-Bench 4.0: 26% vs xAI's claimed 38%; rivals at 60% (GPT-6 Astra) and 55% (Fable 5.1). Cursor's own September eval: 43.9% vs xAI's claimed 46.3%. — [Artificial Analysis \(independent\)](#)

What it means

- Genuine bright spots: Harvey Legal 19.6% leads the field; AA-Briefcase 1657 Elo is top-tier alongside Opus 5 and Fable 5.1 (+111 over 4.6); GDPval-AA 1695 (+90). — [Artificial Analysis \(independent\)](#)
- Average task cost \$4.69 vs ~\$9 for rivals — xAI is competing on perf-per-dollar even where it's mid-pack on capability. — [Cursor \(independent eval\)](#)
- Launch context: five slipped release dates; Grok 4.8 (2.5T params) announced before 4.7 shipped; Musk pre-downgraded expectations Sept 14 to 'roughly on par with Opus 5.0, not 5.1'. — [Musk / xAI](#)

Sources: [xAI / AI Release Tracker](#) · [xAI \(Musk's figure; no spec sheet\)](#) · [AI Release Tracker](#) · [xAI \(company claim\)](#) · [xAI launch materials](#) · [xAI launch materials](#) · [Artificial Analysis \(independent\)](#) · [Artificial Analysis \(independent\)](#) · [Artificial Analysis \(independent\)](#) · [Cursor \(independent eval\)](#) · [Musk / xAI](#)

Compiled from X posts and traditional news outlets. Verify before acting on any single report.